

Computing by Robust Transience: How the Fronto-Parietal Network Performs Sequential, Category-Based Decisions

Highlights

- Recurrent networks trained to perform DMC tasks exhibit robust transience dynamics
- Dynamics consist of stable and slow states connected by robust trajectory tunnels
- Models' neural activities are remarkably similar to recordings from LIP and PFC
- Trained RNNs replicate categorization studies with multiple categories

Authors

Warasinee Chaisangmongkon,
Sruthi K. Swaminathan,
David J. Freedman, Xiao-Jing Wang

Correspondence

xjwang@nyu.edu

In Brief

Chaisangmongkon et al. present a recurrent neural network model of primate fronto-parietal network that can capture various phenomena from neurophysiological experiments in delayed match-to-category tasks.



Computing by Robust Transience: How the Fronto-Parietal Network Performs Sequential, Category-Based Decisions

Warasinee Chaisangmongkon,^{1,2} Sruthi K. Swaminathan,³ David J. Freedman,^{3,4} and Xiao-Jing Wang^{1,5,6,7,*}

¹Department of Neurobiology and Kavli Institute for Neuroscience, Yale University School of Medicine, New Haven, CT 06511, USA

²Institute of Field Robotics, King Mongkut's University of Technology Thonburi, Bangkok 10140, Thailand

³Department of Neurobiology, The University of Chicago, Chicago, IL 60637, USA

⁴Grossman Institute for Neuroscience, Quantitative Biology, and Human Behavior, Chicago, IL 60637, USA

⁵Center for Neural Science, New York University, New York, NY 10003, USA

⁶NYU-ECNU Joint Institute of Brain and Cognitive Science, NYU-Shanghai, Shanghai 200122, China

⁷Lead Contact

*Correspondence: xjwang@nyu.edu

<http://dx.doi.org/10.1016/j.neuron.2017.03.002>

SUMMARY

Decision making involves dynamic interplay between internal judgements and external perception, which has been investigated in delayed match-to-category (DMC) experiments. Our analysis of neural recordings shows that, during DMC tasks, LIP and PFC neurons demonstrate mixed, time-varying, and heterogeneous selectivity, but previous theoretical work has not established the link between these neural characteristics and population-level computations. We trained a recurrent network model to perform DMC tasks and found that the model can remarkably reproduce key features of neuronal selectivity at the single-neuron and population levels. Analysis of the trained networks elucidates that robust transient trajectories of the neural population are the key driver of sequential categorical decisions. The directions of trajectories are governed by network self-organized connectivity, defining a “neural landscape” consisting of a task-tailored arrangement of slow states and dynamical tunnels. With this model, we can identify functionally relevant circuit motifs and generalize the framework to solve other categorization tasks.

INTRODUCTION

Many human behaviors can be viewed as a series of category-based computations (Roelfsema et al., 2003; Rabinovich and Varona, 2011). For example, shopping requires determining categories of desired items, then using that category information to guide our navigation of the store. Neurophysiological studies have investigated the neural basis of such behavior using delayed match-to-category (DMC) tasks in which monkeys indicate whether a test stimulus is a categorical match to a previously presented sample stimulus. Previous work revealed that, in animals trained to perform DMC tasks, neural responses in the

lateral intraparietal area (LIP) and prefrontal cortex (PFC) can exhibit robust selectivity for the category of the sample stimulus that persists across the memory delay period (Freedman et al., 2001; Freedman and Assad, 2006; Swaminathan and Freedman, 2012).

Few computational models have addressed the neural mechanisms underlying sequential category computations as in the DMC task. One class of models suggests that serial, categorical decisions rely on an interplay between multiple subpopulations that each encode specific task parameters, such as stimulus features, categories, rules, and choices (Amit et al., 1994; Ardid and Wang, 2013; Engel and Wang, 2011). These subpopulations are endowed with strong mutual excitations among neurons that prefer the same stimulus feature, driving a stable memory of that feature (Wong and Wang, 2006; Wang, 2001). These models elucidated candidate mechanisms of delay-period-persistent selectivity, but they are not designed to address the diversity or temporal variability in neural responses that are apparent in neural data. In particular, the majority of neurons are responsive to multiple task variables (Raposo et al., 2014; Rigotti et al., 2013; Ibos and Freedman, 2014, 2016; Mante et al., 2013; Freedman and Assad, 2016), and neural encoding often shows baffling temporal variability (Brody et al., 2003a; Crowe et al., 2010; Jun et al., 2010; Shafi et al., 2007). To address these phenomena, recent studies considered an alternative hypothesis that information is distributed randomly within the neural network (Rigotti et al., 2010; Raposo et al., 2014). The idea can be implemented with a network with random connectivity, and to generate different behaviors, downstream circuits can read out relevant information through optimized synaptic weights (Jaeger, 2001; Maass et al., 2002; Rigotti et al., 2010). However, random networks generally do not capture task-specific representations, which can only be acquired through learning. In this realm, we lack a unified framework that can recapitulate all these diverse experimental findings.

To address this problem, we trained a recurrent network model to solve DMC tasks and compared the dynamics of the model network to LIP and PFC data from monkeys performing the DMC task (Freedman and Assad, 2006; Swaminathan and Freedman, 2012). We found that appropriately trained networks

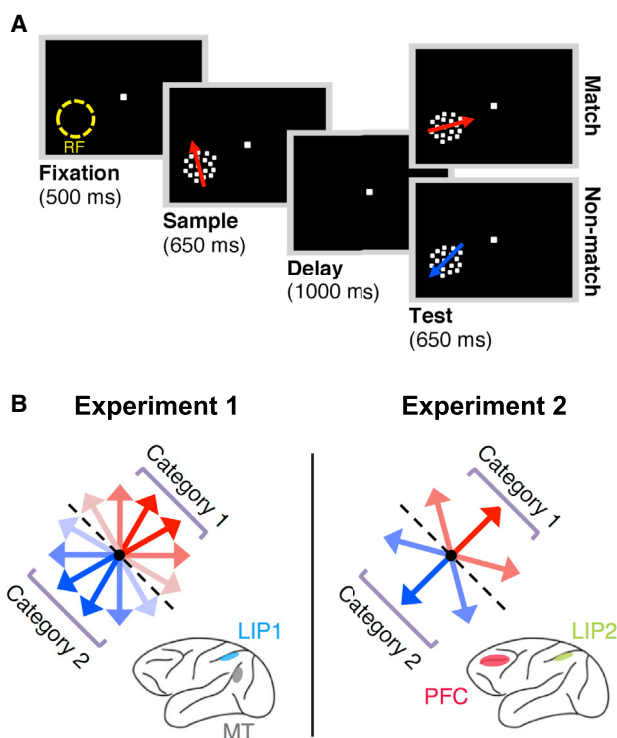


Figure 1. Delayed Match-to-Category Task and Neurophysiological Recordings

(A) Time course of a delayed match-to-category (DMC) experiment. A sample stimulus is followed by a short delay and a test stimulus. To receive reward, subjects must respond whether the sample and test stimuli belong to the same (match) or different (non-match) categories.

(B) Sample and test stimuli are randomly drawn from a set of dot-motion stimuli divided into two categories (red and blue arrows). We analyzed neural recordings from two experiments. The first experiment used 12 motion directions, and LIP neurons were recorded (denoted LIP1). The second experiment used 6 motion directions, and neurons from LIP (denoted LIP2) and PFC were recorded.

reproduce key features of category-dependent responses in the neural data that are not accounted for by previous models.

RESULTS

Delayed Match-to-Category Task and Model Architecture

We analyzed neural recordings from the studies of Freedman and Assad (2006) and Swaminathan and Freedman (2012), in which Macaque monkeys were trained on DMC tasks. In each DMC trial, the stimulus sequence consists of a fixation spot, a sample stimulus, a delay period, and a test stimulus (Figure 1A). Both sample and test stimuli were randomly drawn from a set of random-dot-motion directions that were evenly spaced from 0° to 360° and divided arbitrarily into two categories (marked by red and blue colors in Figure 1B). Subjects learned to report whether the sample and test stimuli belong to the same category (match) or different categories (non-match). The first and second study used 12 and 6 evenly spaced motion directions (30° and

60° apart), respectively. For the first study, we have 156 lateral intraparietal (LIP) neurons from two monkeys. For the second study, we have 74 LIP and 380 prefrontal neurons (PFC) in two other monkeys. We refer to the LIP populations from the first and second experiments as the LIP1 and LIP2, respectively.

Previous studies showed that firing rates of LIP and PFC neurons are markedly tuned to the learned stimulus categories, although the strengths and latencies of categorical signals may differ across areas (Freedman and Assad, 2006; Swaminathan and Freedman, 2012; Swaminathan et al., 2013). The broad similarity in category-related responses suggests that they play overlapping roles in solving the DMC task (Goodwin et al., 2012; Merchant et al., 2011). Instead of stressing the difference between the two regions, this work focuses on understanding the common response patterns observed in both areas.

We trained a recurrent neural network to solve the DMC task. The recurrent network represents a cortical micro-circuit in either the prefrontal or the parietal region; this micro-circuit receives sensory information from visual areas and sends signals to trigger movements in motor areas (Andersen et al., 1990; Cromer et al., 2011; Lewis and Van Essen, 2000; Miller et al., 2002). The network is sparsely connected to noisy input neurons that encode the direction of the sample and test stimuli, mimicking direction-tuned activity in area MT (Figure 2B, top; Freedman and Assad, 2006; Born and Bradley, 2005). A subset of the recurrent population is connected to two action neuron pools, whose activities reflect match or non-match decisions. All connections (input, output, and recurrent) are trained with a supervised method (a Hessian-free algorithm; Martens and Sutskever, 2011; Mante et al., 2013), which adjusts synaptic weights to minimize the difference between the network outputs and specified target responses (i.e., to minimize errors). We instructed the match neuron to hold activity at zero from the beginning of the trial through the delay period, then either to reach a value of 5 (in arbitrary units) at 200 ms after test stimulus onset on match trials or to remain at zero throughout non-match trials. The analogous pattern holds for the non-match neurons. To determine the model's choice, the action neurons' activity is passed through a nonlinear threshold function (Figure 2B, bottom; see also STAR Methods). The match (or non-match) choice is selected when the function value of the match (or non-match) neuron is higher than a threshold of 0.85. We added other output neurons to help stabilize the networks during the resting and post-choice periods (see Figures S1A and S1B).

There are many network configurations that can produce the appropriate output given the specified sensory input (Mante et al., 2013; Barak et al., 2013; Sussillo, 2014). We guide the algorithm to find a subset of solutions that comply with biological constraints by employing additional training strategies (see Figure S1 and STAR Methods). First, the activity of recurrent neurons is restricted to positive values, as is true for neuronal firing rates. Second, the network is trained not only to minimize errors but also to attain sparse synaptic connections. This is achieved by constraining the norm of synaptic weights and eliminating weak synapses iteratively. The target probability of connection is approximately 12%, which is comparable to measurements from mammalian cortical circuits (Song et al., 2005). Third, single neurons should exhibit low spontaneous firing rates when the

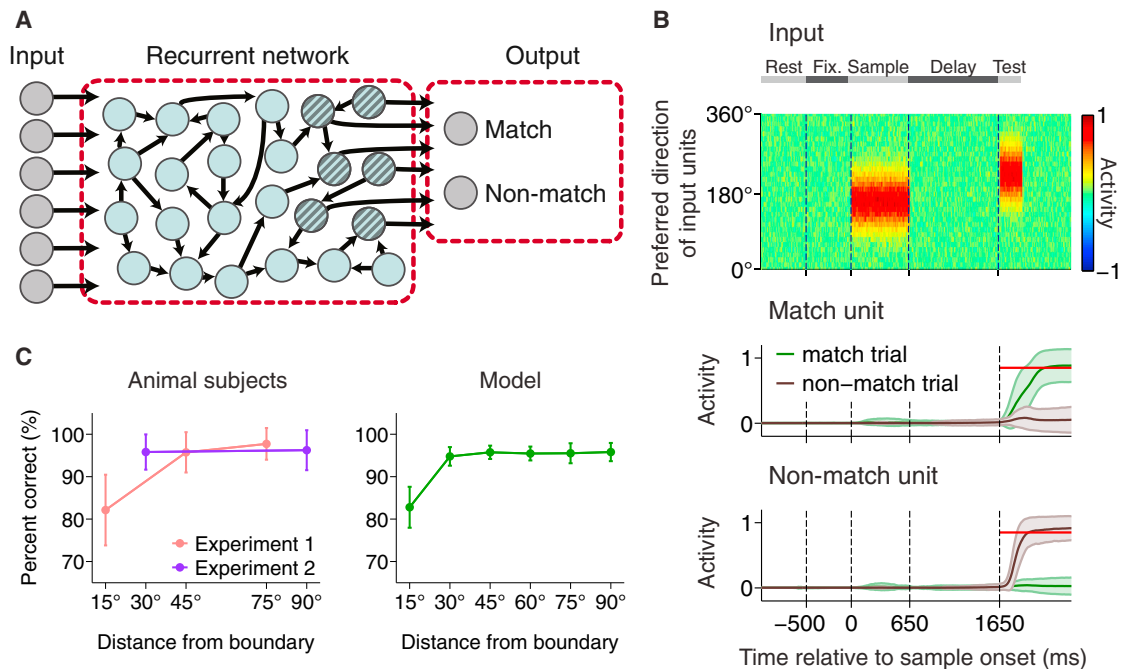


Figure 2. Model Structure and Training Protocol

(A) A set of recurrently connected neurons were trained to solve the DMC paradigm. The recurrent population is connected to direction-selective input units, and approximately one-fifth of recurrent neurons (blue hatched circles) are connected to two output units representing match and non-match choices. All synapses are updated with a supervised learning algorithm.

(B) Activity of the input neurons encode directions of sample and test stimuli (top panel). *x*-axis, time; *y*-axis, input neurons labeled by preferred directions. Neural activity is color-coded. After training, the model generates appropriate decision output. The match output unit (center) shows ramping activation during match trials (green trace) and remains silent during non-match trials (brown trace). The opposite pattern holds in the non-match output unit (bottom). *x*-axis, time; *y*-axis, neural activity; shaded areas, SD across trials. Red lines mark the activity threshold where behavioral choice is registered.

(C) The psychometric functions of animal subjects (left) compared to that of model networks (right). Error bars indicate SD across all recording sessions for animals and across ten network realizations for model.

network is not performing the task. To satisfy this requirement, we instructed the network to hold the sum of all neurons' activity to a small value for 1 s before the trial onset (mean activity = 0.01). Fourth, we employed a progressive training protocol similar to that used to train monkeys on the DMC task (Freedman and Assad, 2006; Swaminathan and Freedman, 2012), whereby the network first started by learning the easiest version of the DMC task with only two stimuli; then, intrinsic noise and more stimuli are gradually introduced. Lastly, the delay durations varied slightly from trial to trial during training (0.9–1.1 ms), resulting in a model that can perform the task with a larger range of delays (Barak et al., 2013; Figure S1C). These constraints and modifications greatly enhanced success rate and the quality of training outcomes.

Training was terminated when the accuracy of the model matched the average performance of animal subjects (88.76%). Both the model and monkeys are less accurate when categorizing near-boundary stimuli (e.g., 15° away from the boundary; Figure 2C).

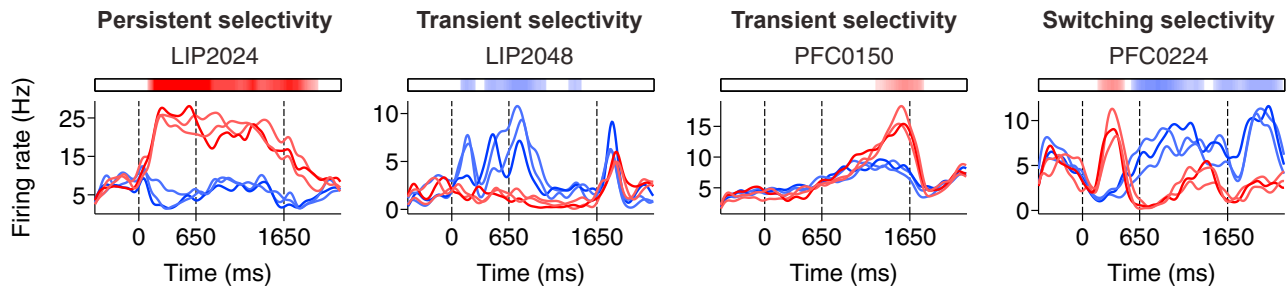
The resulting trained networks provide a candidate dynamical mechanism for solving DMC tasks. To test whether the trained network uses mechanisms similar to real neuronal networks, the model activity must be compared to the recorded experimental neural data.

Heterogeneity in Temporal Profiles of Category Selectivity

We compared the temporal profiles of category selectivity of LIP and PFC neural recordings to those of trained model networks. To quantify the temporal properties of category selectivity, we tested whether the firing rates at each time window are significantly modulated by stimulus categories (*t* tests, $p < 0.05$, Bonferroni corrected). A "category selectivity phase" is defined as a series of consecutive time windows where neural activity shows significant modulation by stimulus categories. The neuron's category selectivity duration is the duration of its longest selectivity phase. The strength of category encoding for each time bin within selectivity phases was quantified by the sensitivity index or d' (see STAR Methods).

Figure 3A illustrates variability in the temporal profiles of category-dependent firing rates for LIP and PFC neurons. Many neurons showed persistent category-dependent responses during the delay period (Figure 3A, far left). In the majority of neurons, the category-selective firing rates undergo marked changes through the delay period. For instance, the category-dependent firing pattern may decay before the end of the delay (Figure 3A, center left) or may commence in the middle of delay (Figure 3A, center right). Furthermore, some neurons switch their category preference in the middle of a trial (Figure 3A, far right). In

A LIP and PFC activity



B Model activity

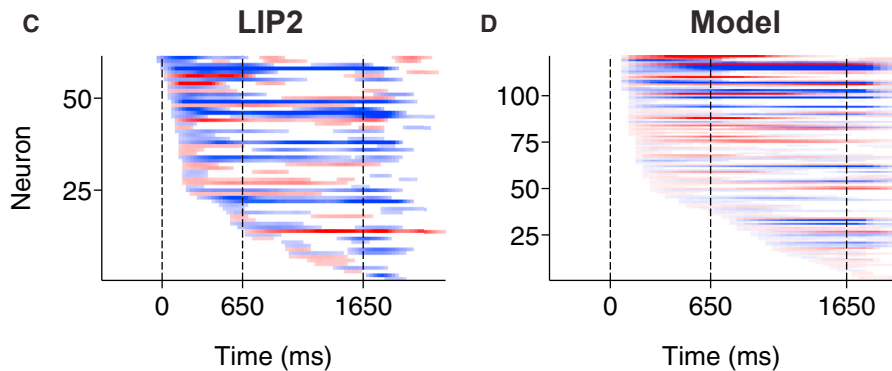



Figure 3. Both Neural Recordings and Model Networks Demonstrate Heterogeneity in the Temporal Profiles of Category Selectivity

(A) Examples of different classes of category selectivity profiles from LIP and PFC populations. Average firing response as a function of time, color-coded by stimulus directions from red to blue category. *x*-axis, time from sample onset; *y*-axis, average firing rate. The colored bar on top shows category-selective period. Color (red or blue) indicates neurons' category preference, while color intensity indicates category tuning strength.

(B) Average neural response of model units with same plotting convention as in (A). Overall, both neural data and model demonstrate heterogeneity in selectivity time course, such as persistent, transient, and switching selectivity.

(C and D) Colored heatmaps showing category selectivity profiles of all neurons in the LIP2 dataset ($N = 61$, C) and of a trained model ($N = 122$, D) with same color-coding convention as colored bars in (A) and (B). Neurons with no selectivity phase are excluded. Across both neural data and model, neurons' category selectivity latencies and durations are highly variable. *x*-axis, time; *y*-axis, neurons sorted by category selectivity latency.

congruence to the neural data, the trained model exhibits heterogeneity in category selectivity profiles, in which persistent, transient, and switching selectivity patterns are observed (Figure 3B). We also quantified fractions of neurons with persistent, transient, and switching selectivity and found that neural data and model show similar trends (Figure S2A).

To visualize the heterogeneity at the population level, we plotted d' of all neurons sorted by the onset of category selectivity from earliest to latest (Figures 3C and 3D). All heatmaps of neural data confirmed two important observations (LIP2 dataset, Figure 3C; LIP1 and PFC datasets, Figures S2B and S2C).

First, neurons can become selective to categories at any time point during the trial. The category selectivity phase does not necessarily align with or overlap with the sample stimulus, which originates the category memory. Second, we observed heterogeneity in the duration of category selectivity across the population. The distribution of category selectivity durations shows a long tail; most neurons exhibit selectivity over short durations, but a small fraction of persistent neurons are consistently detected in all datasets (Figures S2D–S2G). The pattern of category selectivity in model populations reproduces all main features of the neural data.




Figure 4. The Model Captures Essential Features in Population Response Patterns of LIP and PFC Neurons

(A–C) Neural response trajectories during sample period from LIP1 (A), LIP2 (B), and PFC (C) datasets. During the sample period (left column), trajectories begin at roughly the same location for all stimulus directions (black dots correspond to the onset of sample epoch), then fan out into elliptical shapes encoding the directions of stimuli and some category information. Colors of traces encode stimulus identities by the same convention as in Figure 3.

(E–G) Neural response trajectories during delay period from LIP1 (E), LIP2 (F), and PFC (G). During the delay (middle columns), population trajectories have two main components encoding stimulus categories in stable and time-varying manners, respectively. x-axis, relative rate (principal component score); y-axis, time relative to sample onset.

(I–K) Neural response trajectories during test period from LIP1 (I), LIP2 (J), and PFC (K). At the beginning of test period (right column), trajectories encode sample categories (black dots), then neural traces diverge into four separate clusters encoding the four possible sample-test category combinations. BB (dark blue lines) corresponds to blue sample category and blue test category

condition; similarly, RR (dark red) corresponds to red sample and red test, BR (purple) corresponds to blue sample and red test, and RB (orange) corresponds to red sample and blue test. Finally, the traces for match conditions (BB and RR) unfold along the same direction; analogously, so do non-match conditions (BR and RB).

(D, H, and L) Population trajectories of a representative model instance analyzed by the same procedures. The model reproduces population response patterns of neural data in sample (D), delay (H), and test (L) epochs.

Furthermore, our trained networks reproduce mixed category and match selectivity, which is evident in our neural data (Figures S2H and S2I) and other studies (Ibos and Freedman, 2014, 2016; Park et al., 2014; Rishel et al., 2013; Mante et al., 2013; Rigotti et al., 2013). This suggests that a large portion of cortical neurons and model neurons participate in more than one computation.

Population Response Trajectories

We investigated whether neural data and the model exhibit similar patterns of population response trajectories. To this end, we first analyzed the dynamics of neural population responses in LIP and PFC. A neural state is a point in high-dimensional state space where each dimension corresponds to the average firing rate of a neuron at a given time. As neural activity changes over time, a sequence of neural states at consecutive time points forms a population trajectory through state space. For each population, we visualized neural trajectories in a low-dimensional subspace that is most responsive to task conditions using demixing principal component analysis (DPCA) (Machens, 2010; Brendel et al., 2011; Machens et al., 2010), which finds a small set of orthogonal axes that not only captures the most variance in data (like standard PCA) but also segregates response variability due to different task variables onto separate axes. In our case, DPCA yields population response components ranging from one that captures the most variance due to task conditions to one that captures the most variance due to changes in time. We applied DPCA to the mean population response during the

sample (–100–650 ms relative to sample onset), delay (800–1,550 ms), and test (1,600–2,150 ms) epochs separately (see STAR Methods). Note that task conditions were defined by sample motion directions during the sample and delay periods and by sample categories and test directions during the test period.

For the sample period, we applied DPCA, removed the most time-sensitive component, and represented the remaining components on a 2D axis by a multidimensional scaling analysis (Figures 4A–4C; see STAR Methods). At the beginning of the sample epoch, population trajectories originated at the same baseline for all stimuli (black dots, Figures 4A–4C); they then fan out radially, discriminating different sample directions. At the end of stimulus presentation (colored dots, Figures 4A–4C), neural states for all stimulus directions appear in an elliptical configuration, where the stimuli at the middle of both categories (dark blue and red dots, Figures 4A–4C) elicit more distinct population responses than the stimuli close to category boundaries (light blue and red dots, Figures 4A–4C). Overall, LIP and PFC populations show a mixed encoding of sample directions and categories, consistent with our earlier report in Engel et al. (2015).

For the delay epoch, we applied DPCA to delay responses (800–1,550 ms) and projected responses of the whole trial (–250–1,900 ms) onto DPCA axes (see STAR Methods). Here we show two components that participate in the maintenance of categorical working memory (Figures 4E–4G; see also Figure S3). Components in the first column capture the most

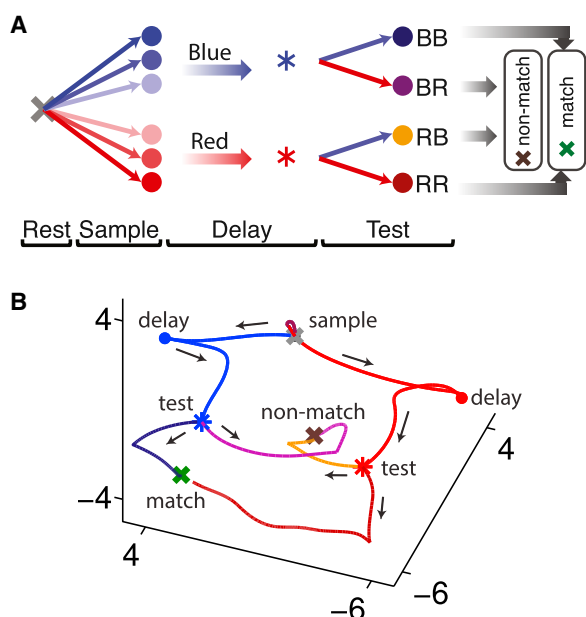


Figure 5. Overall Dynamical Landscape of the Trained Network

(A) A conceptual schematic portraying the series of transitions between behavioral epochs to solve the DMC task.

(B) Neural trajectories of the model implementing the computational process in (A). Neural states evolve serially from the resting state (gray cross) to states associated with sample categories (red and blue dots), then category working memory (red and blue stars), sample-test category combinations (dark blue, dark red, orange, and purple lines), and finally match and non-match choices (green and brown crosses, respectively). Task epoch labels (sample, delay, test) indicate neural states at the beginning of the epoch. All crosses denote stable states; stars denote slow or fixed points associated to working memory, and dots denote transient states.

variance in the delay response due to changes in stimuli (35.5%, 66.3%, and 53.2% for LIP1, LIP2, and PFC, respectively), which constitute a much larger proportion than the second largest component (4.9%, 4.0%, and 10.2%). The components in the first column depict strong and stable encoding of sample categories, representing the main mode of working memory (Figures 4E–4G, left side). The second column shows components with the largest mixture of variance due to changes in time and stimuli, and the neural traces show time-varying category working memory that switches categorical preference mid-delay (Figures 4E–4G, right side). Notably, neural trajectories during the sample and delay epochs show that LIP and PFC populations encode categories by several independent components with different temporal profiles.

For the test epoch, the procedure was similar to the sample epoch except that the neural response was averaged across trials that share the same sample category and test direction. The neural trajectories are grouped into four conditions: BB (dark blue color, Figures 4I–4K) corresponds to trials with blue sample category and blue test category; RR (dark red, Figures 4I–4K) corresponds to red sample and red test; BR (purple, Figures 4I–4K) corresponds to blue sample and red test; and RB (orange, Figures 4I–4K) corresponds to red sample and blue test. At the

beginning of the test period, neural trajectories are clustered according to the sample categories (black dots, Figures 4I–4K). As the test period evolves, the neural traces diverge into four separate clusters encoding sample and test category combinations (BB, RR, BR, and RB conditions). Finally, the trajectories corresponding to match conditions (BB and RR) travel toward the same location, and analogously so, for non-match conditions (BR and RB). Overall, the test-related trajectories encode sample-test category combinations and form states corresponding to match and non-match decisions.

The population trajectories of the trained model networks remarkably reproduce the main features of those from LIP and PFC for all task epochs when the same analyses are applied (Figures 4D, 4H, and 4L). Mixed-direction encoding and categorical encoding are apparent during the sample period (Figure 4D). The population dynamics during the delay encode categories in both stable and time-varying manners (Figure 4H). Lastly, population responses encode sample-test category conditions and converge toward match or non-match states during the test period (Figure 4L). Note that, although the model incorporates a large amount of noise and heterogeneity, the neural data tend to show more variability, especially variance due to changes in time within the trial, which may reflect a timing signal not incorporated in our model (Figure S3). Furthermore, the model tends to be more category-selective than the data, perhaps because the model is a smaller network that is exclusively trained on DMC tasks.

Robust Transient Dynamics Underlie DMC Computations

We characterized the dynamical mechanisms of the model, focusing on two objectives. First, the existence and abundance of mixed-selectivity neurons (Ibos and Freedman, 2014, 2016; Mante et al., 2013; Park et al., 2014; Raposo et al., 2014; Rigotti et al., 2013; Rishel et al., 2013) and the time-varying selectivity for task-relevant variables at both single-neuron (Brody et al., 2003a; Jun et al., 2010; Shafi et al., 2007) and population levels (Machens, 2010; Crowe et al., 2010; Meyers et al., 2008; Wöhrer et al., 2013) have sparked a debate on the dynamical nature of working memory (Druckmann and Chklovskii, 2012; Goldman, 2009; Sarma et al., 2016; Savin and Triesch, 2014; Singh and EliaSmith, 2006). Since the model captures all these features, it is possible now to pinpoint the working memory dynamics that give rise to these patterns of selectivity. Second, we sought to understand how the sequential categorical computations are carried out—i.e., how the sample category information is encoded, maintained, and combined with the test category to generate appropriate behavioral choices.

Figures 5A–5B illustrates the overall trajectories of the model network during DMC tasks, visualized by plotting the largest three principal components of the neural activity. The state space contains key attracting fixed points or unstable saddle points (at which neural activity has near-zero velocity). While performing the task, the network undertakes slow and reliable transitions through these key regions. The directions of movement are determined by the current state location and input. Neural states evolve from the resting state (gray cross, Figures 5A and 5B) to states associated with sample categories (red and blue

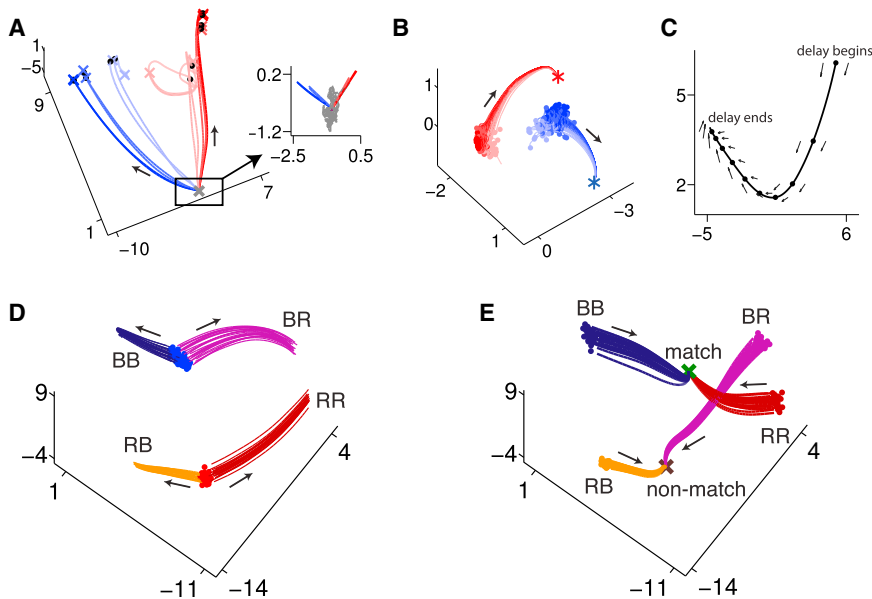


Figure 6. The Trained Network Forms a Dynamical Landscape that Gives Rise to Robust Trajectories and Executes Category-Based Computations

(A) Neural trajectories during sample categorization. At the beginning of a trial, neural states stay within the basin of attraction of the resting state fixed point (gray cross), even upon receiving prolonged (1 s) noisy input (inset, three noisy trajectories plotted in gray). Under the influence of direction-tuned inputs (due to sample stimulus presentation), a set of stable fixed points appear in state space (red and blue crosses, colors denote stimulus directions), propelling the states toward areas associated to red or blue categories.

(B and C) Neural landscape associated to category working memory. (B) Red and blue dots mark possible locations of the neural states at the end of the sample period (with noise). Red and blue lines show noiseless trajectories during the delay originating from these positions. Black arrows mark flow directions. (C) The line shows an example neural state path selected from the trajectories in (B). Black dots mark the neural states at different time points during the delay in 150 ms increments. Arrows show a velocity vector field at

states nearby the trajectory (with norms scaled down for clarity). The working memory landscape resembles tunnels that force neural states to flow along two possible routes (arrows in C show the movement flow within the tunnel), generating robust time-varying memory of sample categories. The neural state movements slow down or stop at the end of delay (marked by shorter distance between dots near the end of delay in C).

(D and E) Neural trajectories during match decisions. (D) Red and blue dots mark locations that neural states occupy at the end of the delay (with noise). When the test stimuli appear, neural states move toward the same input-dependent fixed points as shown in (A). Since there are two possible starting regions (associated to red or blue sample categories), trajectories diverge along four separate streams, encoding sample-test category conditions. (E) Continuing from (D), when test stimuli are removed, neural states fall into the basins of attraction corresponding to match or non-match stable states.

Data plotted come from a representative model instance.

dots mark the end of sample period, [Figures 5A and 5B](#)), then category working memory (red and blue stars mark the end of delay period, [Figures 5A and 5B](#)), sample-test category combinations (dark blue, dark red, orange, and purple dots and lines in [Figures 5A and 5B](#), respectively), and, finally, match or non-match decision states (green and brown crosses, [Figures 5A and 5B](#)). The whole series of state transitions solves the DMC task.

We characterized the network dynamics at each time epoch in more detail. During rest and fixation periods, the network trajectories are confined within the resting state's basin of attraction despite the noisy sensory signals the network receives (gray traces, [Figure 6A inset](#)). When the system receives a stimulus-selective input, stimuli in different categories propel neural states to separate directions ([Figure S4A](#)) and, over time, out of the resting state basin ([Figures 6A and S4B](#)). If the network inputs stay on for a long period, the network would converge to stimulus-dependent steady states (red and blue crosses, [Figure 6A](#)), which are clustered based on stimulus categories. The arrangement of input-dependent stable states ([Rabinovich et al., 2001, 2008](#)) results in directional neural trajectories that distinguish between stimulus categories.

At the end of the sample period, the network's slow and transient states are distributed within two regions in state space (red and blue dots, [Figure 6B](#); variability in locations is a result of noise, and network's low velocity is shown in [Figure S4C](#)). Using these states as initial conditions, we simulated network activity

without input and noise. We observed that network states relax along two narrowing tunnels, one for each sample category, maintaining category memories in a dynamic manner ([Figure 6B](#)). The velocity vector field near one of the tunnel centers is plotted in [Figure 6C](#), where arrow lengths indicate relative velocity magnitude. The plot shows that the neural state moves more slowly as it approaches the end of the tunnel (see also [Figure S4D](#)), and arrow directions point toward the middle of the tunnel, funneling the system's state to a specific region near the end of the delay. This dynamical analysis revealed that categorical working memory is maintained by robust trajectories, which explains why we observed time-varying selectivity at both single-neuron ([Figure 3](#)) and population ([Figure 4](#)) levels.

Note that states associated with stimuli near the category boundary are closer to each other at the beginning of the delay than states of stimuli further away from the boundary (pale red and blue lines in [Figures 6A and 6B](#)). Misclassification occurs when the end-of-sample states stray outside of the tunnel under the influence of noise and end up either in the wrong categorical tunnel or in the basins associated to rest state or choices ([Figure S4E](#)). This takes place more often for near-boundary stimuli, leading to poorer performance, as shown in [Figure 2C](#).

The end-of-delay regions are in the proximity of saddle or stable points (stability of this state varies across networks trained by an identical protocol), leading to low network velocity and keeping the memory of categories for an extended period. This allows the networks to perform well even when delay durations vary

(accuracy >70% in the range of 0.7–1.3 s delay) (Figure S1C). Note that if the delay is prolonged much longer than 1 s, our simulation shows two possible outcomes. First, if stable states associated to categorical working memory emerged during training, neural states simply rest in stable states. We observe that category-related fixed points are likely to emerge if the network is trained with variable delay duration randomly drawn from a larger range (0.8–2 s, Figure S5). Second, prolonged delays may lead to a gradual decay of working memory whereby the network collapses to fixed points associated to resting state or random choices. The neural datasets we investigate cannot distinguish between these two scenarios.

At the onset of the test period, neural states are distributed between two regions of state space associated to red and blue categories of working memory (red and blue dots in Figure 6D); variability in state locations is due to noise. When the test stimulus is introduced, the direction-selective input shifts the landscape in the same fashion as in the sample period (Figure 6A), directing the network toward stimulus-dependent stable states. However, since trajectories are launched from two possible initial locations depending on the sample category, the neural paths split into four separate streams encoding sample-test category combinations (dark blue, dark red, purple, and orange lines in Figure 6D), bringing neural states to four separate clusters of transient states. This dynamical picture provides a concrete example of state-dependent computations in which the same stimulus can be interpreted differently or lead to different behavioral outcomes depending on the prior experience of the network (Buonomano and Maass, 2009).

Soon after the test stimulus appears, the match (or non-match) output neuron can read out from the recurrent neural states and ramp up to response threshold during the match (or non-match) trial. The response time is usually within a few hundred milliseconds after the test stimulus onset. Finally, after the response is committed and test stimulus is removed, the network relaxes along its natural landscape. The four regions in state space (dots in Figure 6E) are mapped onto two steady states (crosses in Figure 6E); RR and BB traces go to one point (match attractor), while RB and BR traces go to a separate point (non-match attractor). These stable states complete the sequence of DMC computations.

The dynamics of the model reveal that a single cortical network can carry out a series of computations by utilizing different regions of state space to perform different computations. This idea is well supported by recent work investigating sensory encoding (Rabinovich et al., 2001, 2008), decision processes (Raposo et al., 2014; Mante et al., 2013; Murakami and Mainen, 2015), and movement execution (Churchland et al., 2012; Hennequin et al., 2014). To understand computational mechanisms, one must consider the population dynamics as a whole. Observing this network-level phenomenon through the activity of a single neuron amounts to watching a moving object in three-dimensional space through its one-dimensional projection. The projected image may miss salient information (such as categorical discrimination) at some moments, or it may contain information from more than one process. Therefore, mixed and time-varying selectivity are expected and observed at the single neuron level.

Structural and Functional Connectivity of Trained Networks

We performed a series of analyses to understand the connectivity structure that governs robust transient dynamics. We compared our trained recurrent networks with randomly connected networks (RCNs), which were previously investigated as a source of mixed time-varying selectivity (Rigotti et al., 2010; Barak et al., 2013). We found that, although RCNs encode a mixture of stimulus- and time-dependent variability, they do not exhibit the self-generated, categorical neuronal coding during the sample and/or delay periods. In particular, in both neural data and networks with trained recurrent connections, the majority of neurons with persistent selectivity are strong categorical discriminators, whereas in RCNs, persistent representation is not category-specific (Figures S6A–S6F). However, units in our recurrent neural network do display mixed selectivity; therefore, this work extends the work of Rigotti et al. (2010) to networks with wiring structures that emerge from training to perform a cognitive task.

In trained networks, the distribution of synaptic connections is sparse and unimodal with mean weight equal to zero (Figure S6G), but exhibits a clear hierarchical structure not present in RCNs. To reveal the hierarchy, we computed degree centrality, defined as the total number of connections each neuron sends (out-degree) and receives (in-degree). All trained connectivity exhibits heavy-tailed degree distributions—i.e., few neurons are connected to large numbers of neighbors acting as network hubs (Figure 7A, mean kurtosis for all ten networks = 18.341, $p = 0.005$). Furthermore, trained networks also exhibit a high correlation between in-degree and out-degree (Figure 7B; Spearman rank correlation, $N = 150$, $\rho = 0.546$, $P < 10^{-7}$), suggesting that hub neurons aggregate information from them and broadcast it to large numbers of neighbors. Neurons with high degree also tend to have larger positive incoming connections (large in-strength, Figure S6H) and larger average activity (Figure S6I) than their low-degree counterparts, which means that these neurons have greater influence on neural state trajectories.

The robust dynamics underlying sequential decisions result from ongoing competition and cooperation among neurons within the circuit. We measured neural response similarity (r_{response}), defined as the covariance between synaptic currents of neural pairs across task conditions (see STAR Methods), and structural coupling, defined as the sum of synaptic connections from neuron i to j and from j to i . The r_{response} is correlated with structural coupling in all task epochs (Pearson correlation, average $r = 0.176$, $P < 10^{-4}$; see statistical test against null models in Figure S6J), suggesting that neurons with similar category or match selectivity tend to have strong positive synaptic couplings, while neurons with opposite encoding have strong negative couplings. This gives rise to competitive dynamics between subpopulations that encode different concepts (Wong and Wang, 2006; Wang, 2002). To further investigate neural couplings, we divided the neural population into four groups based on their noiseless activity at the end of the delay period: (1) neurons that are active when a stimulus belongs to the red category but silent for the blue category (denoted as R group, average 10.2% of population); (2) neurons that are active exclusively for the blue category (B, 14.7%); (3) neurons that are responsive

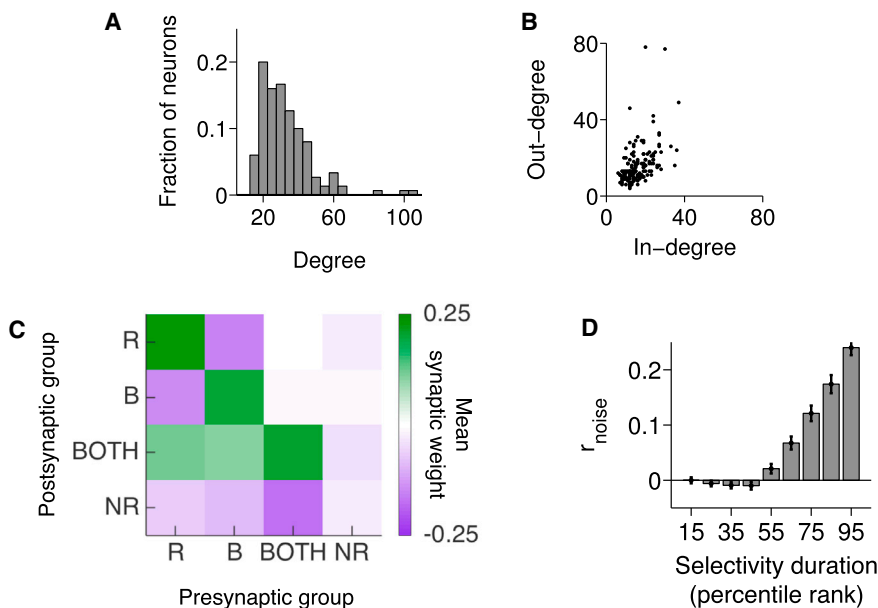


Figure 7. Structural and Functional Connectivity that Supports Robust Transient Dynamics

(A) Degree distribution (total number of connections) in a representative sample of a trained network. All trained networks exhibit long-tail degree distribution, showing existence of hub neurons.

(B) Scatterplot shows the number of incoming connections (in-degree, x axis) versus outgoing connections (out-degree, y axis) for all neural units in a trained network. Strong correlations between in-degree and out-degree are observed in trained networks, which is unexpected if their connectivities are random (Pearson correlation, $N = 150$, $r = 0.482$, $P < 10^{-7}$).

(C) Average synaptic weights between neuron groups that are active only for red stimuli (R group), only for blue (B group), for both red and blue (BOTH group), and not responsive (NR group). Colors indicate average synaptic connection (purple, inhibitory connections; green, excitatory connections) from a pre-synaptic group (x axis) to a post-synaptic group (y axis). R and B groups exhibit within-

group excitation and across-group inhibition, while the BOTH group receives excitatory connections from itself and from R and B groups. The NR group receives inhibitory connections from all groups.

(D) Average noise correlation (r_{noise} , y axis) of neuron pairs grouped by percentile ranks of their average category selectivity durations. The x axis is the center of each rank bin (bin width = 10 percent). Neurons with persistent category-selective activity tend to form large functional connections (high r_{noise}). All ten network realizations demonstrate similar features; data plotted come from a representative sample. Error bars indicate SEM of r_{noise} across neurons in the same bin.

for both red and blue (BOTH, 14.5%); and (4) neurons that are not responsive at all (NR, 60.6%). Then we assessed average connections within and between these subclasses. We found strong within-group excitation and between-group mutual inhibition for R and B groups (Figure 7C), mediating competition between the two categories. Furthermore, we found that neurons in the BOTH group have high degrees and activity but are less sensitive to categories than R and B groups (Figures S7A–S7C). The BOTH group receives net excitatory connections from itself as well as from R and B groups (Figure 7C). The activity of BOTH neurons tends to increase over the delay, while that of R and B neurons tends to decrease (Figure S7D). BOTH neurons' activity drives correlations between neural states associated to red and blue categories, which is apparent in both the neural data and model (Figures S7E and S7F). Overall, these findings show the cooperation between two categories of neural pools through BOTH neurons. This co-activation is likely responsible for the temporal dynamics that bring neural states to the end-of-delay regions, where match and non-match decisions can be made separately from the categorical decision.

Lastly, the model yields a testable prediction that functional and structural coupling among neurons with persistent selectivity that prefer the same category tend to be larger relative to all connections in the network. For a given pair of neurons with the same category preference, we measured the average category selectivity duration (CSD, see STAR Methods) and noise correlation (r_{noise})—i.e., the correlation coefficient between a neuron pair's rate fluctuations averaged across all task conditions. Neural pairs that contain non-selective neurons are

removed from the analysis. We found that persistent neurons in the model tend to have far larger functional couplings than do non-persistent neurons (Figure 7D), whereas neural pairs with average CSD larger than 90 th percentile have a larger average noise correlation ($m_1 = 0.237$) than other pairs ($m_2 = 0.031$, t test, $p < 10^{-7}$). This result holds for any time window at which r_{noise} is measured. The effect remains significant when the same analysis is performed on synaptic coupling instead of on r_{noise} ($m_1 = 0.234$, $m_2 = 0.031$, $p < 10^{-7}$) and when controlled for average neural activity (ANCOVA, $F = 326.69$, $p < 10^{-7}$, see STAR Methods).

Neuronal Representation during Flexible Categorization with Multiple Rules

Recent studies have shown that single neurons in LIP (Fitzgerald et al., 2011) and PFC (Cromer et al., 2010) are multitaskers, as they encode categorical information for different sets of stimuli (e.g., differentiating between dogs versus cats for animal classification task and sports versus sedans in car classification). These studies found that (1) multitasking neurons were the strongest category discriminators (Cromer et al., 2010) and (2) neurons' tuning strengths for different stimulus sets were correlated (Fitzgerald et al., 2011). We refer to these tasks as “independent-input” paradigms, as the two categorical schemes involve independent stimulus sets with likely non-overlapping sensory representations. In contrast, another set of studies employed a different paradigm in which subjects were instructed to categorize the same set of stimuli under two different schemes (e.g., categorizing the same images of animals into dogs versus cats

or big versus small depending on the active rule) (Roy et al., 2010; Goodwin et al., 2012). We refer to these tasks as “shared-input” paradigms because both categorical schemes share the same sensory representation. These experiments show that (1) rule-dependent responses emerged as soon as the rule cue was presented (Goodwin et al., 2012) and (2) multitasking neurons were more commonly observed, whereas specialized neurons (i.e., neurons that encoded categories exclusively for one scheme) were less common in the independent-input paradigm than in the shared-input paradigm (Roy et al., 2010; Cromer et al., 2010).

We asked the following questions: do different task paradigms elicit different dynamical landscapes, and does the discrepancy in dynamical structures alone account for these experimental observations?

To investigate this question, we trained recurrent neural networks, using the same protocol we used for the standard DMC task, to solve either independent-input or shared-input categorization tasks (see STAR Methods). For the independent-input paradigm, one input neuron group encodes motion directions (Figure 8A, scheme A, red and blue categories), while another group encodes spatial locations of a circle stimulus (Figure 8A, scheme B, pink and green categories). The network’s task is to categorize stimuli according to the boundary associated with each stimulus set (dashed black lines in Figure 8A). For the shared-input paradigm, the model learned to categorize motion directions by two different boundaries (Figure 8D). Prior to the fixation epoch, the model received a 500 ms input pulse from two separate input neurons (colored squares in Figure 8D), signifying whether the horizontal (Figure 8D, scheme A) or vertical boundary (Figure 8D, scheme B) is in effect.

The two task paradigms result in markedly different landscapes. In the independent-input case, the trained networks form two working memory tunnels during the delay, similarly to those in Figure 6B, but these tunnels are shared between the two categorical schemes (Figure 8B). In particular, one tunnel corresponds to the red category of scheme A and the green category of scheme B while another tunnel corresponds to the blue and pink categories. Note that the opposite configuration (red/pink and blue/green) is also possible. Only two tunnels are required to solve the independent-input task, since the entrances of the tunnels can be mapped onto appropriate stimuli by modifying separate sets of input weights from different sensory neuron groups (Figure S8A) and the ends of the tunnels are mapped onto appropriate choices by a similar mechanism (Figure S8B). Hence, the two categorical schemes can share the same categorical discrimination machinery via appropriate mapping. Neurons that participate in driving trajectories along the tunnels must be active in both schemes, leading to strong correlation between the category tuning indices (CTIs) of the two schemes (Figure 8C; Pearson correlation, $N = 150$, $r = 0.85$, $P < 10^{-7}$). Furthermore, neurons with persistent contribution to tunnel trajectories tend to be those with the strongest categorical selectivity (Figures S6C and S6D); therefore, multitasking neurons are the most robust category discriminators (Figure 8C; see statistical test in Figure S8C).

Tunnel sharing is only possible when sensory inputs contain rule information—i.e., motion direction stimuli entail a horizontal

boundary and dot location stimuli entail a vertical boundary. In such case, the recurrent network does not need to memorize the categorization rule through the delay to map the content of working memory to the appropriate choices during the test period. This strategy would fail in the shared-input paradigm, where the same set of stimuli must be mapped to different choices during the test period depending on the active rule. Instead, we observed that neural trajectories diverge to encode categorization rules after the rule cue is displayed through the fixation period (Figure 8E), resembling rule representations observed in PFC (Wallis et al., 2001; Goodwin et al., 2012). When the sample stimulus is presented, the dynamic representations encode both rules and categories, which persist through the delay (Figure 8F). Since the two category schemes no longer share the same tunnels, the correlation between category tuning strengths reduces or vanishes (Figure S8D). The overall effects from ten instances of networks trained with either paradigm show that the shared-input paradigm leads to lower correlation in category tuning strengths between the two schemes (Figure 8G; Pearson correlation, $N = 5$ realizations of network; independent input, average $r = 0.76$, $P < 10^{-6}$ in all networks; Shared input, average $r = 0.08$, only one out of five realizations has significant correlation, $P < 0.01$). Consequently, we observed a significantly smaller number of multitasking neurons and a larger number of specialized neurons than in the independent-input paradigm, in accordance with experimental findings (Figure 8H; t test, $N = 5$, $P = 0.001$). Collectively, this comparison between independent-input and shared-input paradigms illustrates how dynamical landscapes can adapt to various categorical structures, and the difference in landscapes alone can explain a lot of experimental findings.

DISCUSSION

Our results contribute four important insights. First, our model suggests that robust transient dynamics, equipped with stimulus-dependent attracting states and robust trajectory tunnels, underlie delayed associative computations in cortical circuits. Second, we show that networks endowed with reproducible trajectories capture statistics of the heterogeneous and time-varying category selectivity at both the single-neuron and population levels, thus bridging the robust transience framework to neurophysiology of the primate fronto-parietal network. Third, we reveal the features of structural and functional connectivity that support robust transience and suggest a testable prediction about the relationship between the temporal profiles of selectivity and inter-neuronal correlations. Fourth, our model explains observations from experiments that incorporate multiple categorization rules through the idea of shared state space landscape.

Much emphasis has been put on the reward-dependent learning mechanism that explains the emergence of categorical representation (Roelfsema and van Ooyen, 2005; Engel et al., 2015; Rombouts et al., 2012; Savin and Triesch, 2014). Though providing valuable insights on synaptic plasticity, many of these studies have not focused on the temporal profiles of category selectivity, and none have evaluated whether the end results of training resemble the neural dynamics in the brain. Our study focuses on the dynamical properties of networks that successfully

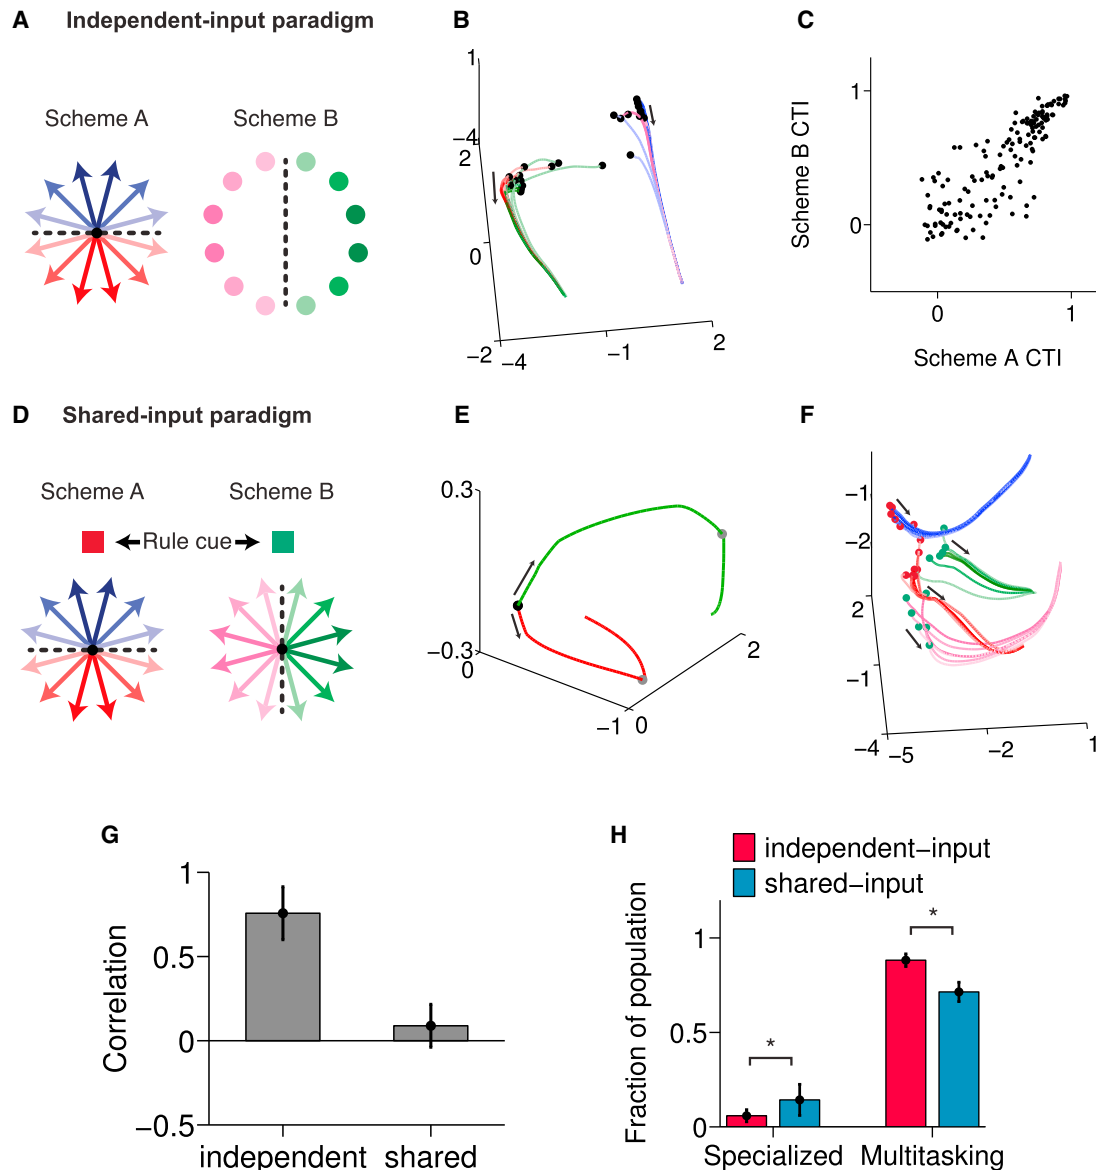


Figure 8. Robust Transience Framework Explains Neural Selectivity during Flexible Categorization involving Multiple Rules

(A) An independent-input categorization paradigm. Networks learn to categorize two separate sets of stimuli (motion directions, scheme A; stationary dots at different spatial locations, scheme B). The two stimulus sets are represented by two separate groups of sensory neurons and are subject to different categorization rules.

(B) Noiseless trajectories during the delay from networks trained with independent-input paradigm (black dots mark the beginning of the delay epoch). Networks form two working memory tunnels that are utilized by both stimulus sets to maintain category working memory.

(C) We calculated neurons' category tuning index (CTI), which measures the strength of categorical sensitivity to any preferred category. Tunnel sharing results in a large correlation between CTIs for scheme A (x axis) versus scheme B (y axis) during the delay period (Pearson correlation, $n = 150$, $r = 0.85$, $P < 10^{-7}$). Data plotted in (B) and (C) come from a representative model instance.

(D) A shared-input categorization paradigm. Networks must categorize a single stimulus set by two different boundaries, signaled by colors of the rule cue.

(E) Noiseless trajectories during the rule cue and fixation periods for a network trained with a shared-input paradigm. Black and gray dots mark the beginnings of the rule cue and the fixation period, respectively. Trajectories split into two streams corresponding to different rules.

(F) Network forms four separate tunnels to maintain category-rule combination. Categorization rules coded by dot colors. Data plotted in (E) and (F) come from a representative model instance.

(G) Correlations between category tuning indices for scheme A and scheme B across five realizations of networks trained with an independent- or shared-input paradigm. Error bars indicate minimum and maximum correlations within each group. Independent-input paradigm results in large positive correlations between CTIs of the two schemes, while shared-input paradigm does not.

(H) The independent-input paradigm produces a significantly smaller number of specialized neurons but a larger number of multitasking neurons than the shared-input paradigm (t test, $n = 5$, stars indicate $P < 0.05$). Error bars indicate SD of fractions across five network realizations.

solve the task and exhibit similar response features to neurophysiological data. The accrued insights provide essential foundations for future generative models. Note that this study assumed that, although neurons were recorded at different times and in different animals, their activities represents sampled firing rates from a single working population.

To the extent possible, our model parameters were calibrated by experimental measurements, such as the sparse connectivity of the trained networks (Song et al., 2005), the neural time constant (Murray et al., 2014), and the width of sensory tuning (Albright, 1984). Other parameters, such as the initial recurrent network connections and noise in the networks, were set in the same range as those in previous modeling studies (Mante et al., 2013). The changes in these parameters do not affect our overall findings, but can impact training. For example, a longer neural time constant will make it easier to train networks on longer delay epochs, and higher noise in the network will reduce the chance of training success.

Importantly, our results suggest that time-varying patterns of category working memory result from a slow dynamic transition from one location in state space to another, mediated by a dynamical tunnel that constrains the course of trajectories. This is distinct from purely feedforward models (Goldman, 2009; Savin and Triesch, 2014) or models that utilize rapid transitions to stable states (Wong and Wang, 2006; Wang, 2002). One accompanying feature of such a mechanism is the reliable emergence of persistently selective neurons among other neurons with heterogeneous temporal dependence. This gives rise to the category-selective population code; its dominant mode is stable, yet it also exhibits a time-varying secondary mode. Similar population dynamics have been observed in other tasks (Machens, 2010; Raposo et al., 2014) but have not been accounted for by other models.

Despite their time-varying dynamics, networks utilizing robust transience support and advance the central idea of strong reverberatory dynamics underlying working memory and decision making (Goldman-Rakic, 1990; Wang, 2001, 2002; Wong and Wang, 2006; Murray et al., 2017). The persistent neurons in our model, albeit few in quantity, are the main drivers of delay dynamics, as they are among the strongest category discriminators and form large connections among one another. The circuit motifs proposed in classical models, such as strong local excitation and mutual inhibition among dominant neuron groups, are apparent in the current framework, although they are embedded in more heterogeneous circuits; this allows them to flexibly partake in sequential computations and to generate mixed representations in accordance with experimental evidence. The presence of multiple stable states is also the key constituent of robust dynamics in our model. Furthermore, structural organization in local circuits may vary in a continuum from random networks to robust dynamics to stable attractors, depending on the extent of training (Barak et al., 2013). In particular, our model predicts that networks trained on protocols in which delay durations vary across trials tend to develop more temporally stable persistent activity (see Figure S5). Future animal experiments can test this hypothesis.

Our work contributes to a growing line of research on robust transient dynamics and their role in complex neural computa-

tions. The principle has been proposed for spatiotemporal sensory encoding (Rabinovich et al., 2001), movement generation (Hennequin et al., 2014), and other cognitive processes (Rabinovich and Varona, 2011; Rabinovich et al., 2008, 2014), which speaks to its prevalence in neural circuit processing across brain regions and species. Most importantly, through detailed comparison between neurophysiological data and model, our contribution provides compelling evidence that robust transience governs sequential categorical decisions in primate cortical circuitry.

STAR★METHODS

Detailed methods are provided in the online version of this paper and include the following:

- KEY RESOURCES TABLE
- CONTACT FOR REAGENT AND RESOURCE SHARING
- EXPERIMENTAL MODEL AND SUBJECT DETAILS
- METHOD DETAILS
 - Model Architecture and Training
 - Model for Multi-scheme Categorization Tasks
 - Analysis of Model Dynamics and Connectivity
- QUANTIFICATION AND STATISTICAL ANALYSIS
 - Analyzing Temporal Properties of Selectivity
 - Population Response Analysis
- DATA AND SOFTWARE AVAILABILITY

SUPPLEMENTAL INFORMATION

Supplemental Information includes eight figures and can be found with this article online at <http://dx.doi.org/10.1016/j.neuron.2017.03.002>.

AUTHOR CONTRIBUTIONS

W.C. and X.-J.W. designed research, performed model simulations, and analyzed data. D.J.F. and S.K.S. designed and performed experiments. W.C., D.J.F., and X.-J.W. wrote the paper.

ACKNOWLEDGMENTS

This work was supported by NIH grants R01MH062349 and R01MH092927, NSF-NCS grant 1631571, and STCSM grants 14JC1404900 and 15JC1400104. We thank John Assad for valuable contributions during all phases of the neurophysiological studies, which produced the data examined here. We thank John Murray, Francis Song, and William Gaines for intellectual and helpful discussions.

Received: May 24, 2016

Revised: September 30, 2016

Accepted: February 27, 2017

Published: March 22, 2017

SUPPORTING CITATIONS

The following references appear in the Supplemental Information: Barak et al. (2010); Bullmore and Sporns (2009); Maslov and Sneppen (2002); Miller et al. (2003); Rubinov and Sporns (2010); Stokes et al. (2013).

REFERENCES

- Albright, T.D. (1984). Direction and orientation selectivity of neurons in visual area MT of the macaque. *J. Neurophysiol.* 52, 1106–1130.
- Amit, D.J., Brunel, N., and Tsodyks, M.V. (1994). Correlations of cortical Hebbian reverberations: theory versus experiment. *J. Neurosci.* 14, 6435–6445.
- Andersen, R.A., Asanuma, C., Essick, G., and Siegel, R.M. (1990). Corticocortical connections of anatomically and physiologically defined subdivisions within the inferior parietal lobule. *J. Comp. Neurol.* 296, 65–113.
- Ardid, S., and Wang, X.-J. (2013). A tweaking principle for executive control: neuronal circuit mechanism for rule-based task switching and conflict resolution. *J. Neurosci.* 33, 19504–19517.
- Barak, O., Tsodyks, M., and Romo, R. (2010). Neuronal population coding of parametric working memory. *J. Neurosci.* 30, 9424–9430.
- Barak, O., Sussillo, D., Romo, R., Tsodyks, M., and Abbott, L.F. (2013). From fixed points to chaos: three models of delayed discrimination. *Prog. Neurobiol.* 103, 214–222.
- Borg, I., and Groenen, P.J.F. (1997). *Modern Multidimensional Scaling: Theory and Applications* (Springer-Verlag).
- Born, R.T., and Bradley, D.C. (2005). Structure and function of visual area MT. *Annu. Rev. Neurosci.* 28, 157–189.
- Brendel, W., Romo, R., and Machens, C.K. (2011). Demixed principal component analysis. In *Advances in Neural Information Processing Systems*, 2654–2662.
- Brody, C.D., Hernández, A., Zainos, A., and Romo, R. (2003a). Timing and neural encoding of somatosensory parametric working memory in macaque prefrontal cortex. *Cereb. Cortex* 13, 1196–1207.
- Bullmore, E., and Sporns, O. (2009). Complex brain networks: graph theoretical analysis of structural and functional systems. *Nat. Rev. Neurosci.* 10, 186–198.
- Buonomano, D.V., and Maass, W. (2009). State-dependent computations: spatiotemporal processing in cortical networks. *Nat. Rev. Neurosci.* 10, 113–125.
- Churchland, M.M., Cunningham, J.P., Kaufman, M.T., Foster, J.D., Nuyujukian, P., Ryu, S.I., and Shenoy, K.V. (2012). Neural population dynamics during reaching. *Nature* 487, 51–56.
- Cohen, M.R., and Kohn, A. (2011). Measuring and interpreting neuronal correlations. *Nat. Neurosci.* 14, 811–819.
- Cromer, J.A., Roy, J.E., and Miller, E.K. (2010). Representation of multiple, independent categories in the primate prefrontal cortex. *Neuron* 66, 796–807.
- Cromer, J.A., Roy, J.E., Buschman, T.J., and Miller, E.K. (2011). Comparison of primate prefrontal and premotor cortex neuronal activity during visual categorization. *J. Cogn. Neurosci.* 23, 3355–3365.
- Crowe, D.A., Averbeck, B.B., and Chafee, M.V. (2010). Rapid sequences of population activity patterns dynamically encode task-critical spatial information in parietal cortex. *J. Neurosci.* 30, 11640–11653.
- Druckmann, S., and Chklovskii, D.B. (2012). Neuronal circuits underlying persistent representations despite time varying activity. *Curr. Biol.* 22, 2095–2103.
- Engel, T.A., and Wang, X.-J. (2011). Same or different? A neural circuit mechanism of similarity-based pattern match decision making. *J. Neurosci.* 31, 6982–6996.
- Engel, T.A., Chaisangmongkon, W., Freedman, D.J., and Wang, X.-J. (2015). Choice-correlated activity fluctuations underlie learning of neuronal category representations. *Nat Commun* 6, 6454.
- Fitzgerald, J.K., Freedman, D.J., and Assad, J.A. (2011). Generalized associative representations in parietal cortex. *Nat. Neurosci.* 14, 1075–1079.
- Freedman, D.J., and Assad, J.A. (2006). Experience-dependent representation of visual categories in parietal cortex. *Nature* 443, 85–88.
- Freedman, D.J., and Assad, J.A. (2016). Neuronal mechanisms of visual categorization: an abstract view on decision making. *Annu. Rev. Neurosci.* 39, 129–147.
- Freedman, D.J., Riesenhuber, M., Poggio, T., and Miller, E.K. (2001). Categorical representation of visual stimuli in the primate prefrontal cortex. *Science* 291, 312–316.
- Goldman, M.S. (2009). Memory without feedback in a neural network. *Neuron* 61, 621–634.
- Goldman-Rakic, P.S. (1990). Cellular and circuit basis of working memory in prefrontal cortex of nonhuman primates. *Prog. Brain Res.* 85, 325–335, discussion 335–336.
- Goodwin, S.J., Blackman, R.K., Sakellari, S., and Chafee, M.V. (2012). Executive control over cognition: stronger and earlier rule-based modulation of spatial category signals in prefrontal cortex relative to parietal cortex. *J. Neurosci.* 32, 3499–3515.
- Hennequin, G., Vogels, T.P., and Gerstner, W. (2014). Optimal control of transient dynamics in balanced networks supports generation of complex movements. *Neuron* 82, 1394–1406.
- Ibos, G., and Freedman, D.J. (2014). Dynamic integration of task-relevant visual features in posterior parietal cortex. *Neuron* 83, 1468–1480.
- Ibos, G., and Freedman, D.J. (2016). Interaction between spatial and feature attention in posterior parietal cortex. *Neuron* 91, 931–943.
- Jaeger, H. (2001). The “echo state” approach to analysing and training recurrent neural networks—with an erratum note. GMD Technical Report 148, German National Research Center for Information Technology. <https://pdfs.semanticscholar.org/8430/c0b9afa478ae660398704b11dca1221ccf22.pdf>.
- Jun, J.K., Miller, P., Hernández, A., Zainos, A., Lemus, L., Brody, C.D., and Romo, R. (2010). Heterogeneous population coding of a short-term memory and decision task. *J. Neurosci.* 30, 916–929.
- Lewis, J.W., and Van Essen, D.C. (2000). Corticocortical connections of visual, sensorimotor, and multimodal processing areas in the parietal lobe of the macaque monkey. *J. Comp. Neurol.* 428, 112–137.
- Maass, W., Natschläger, T., and Markram, H. (2002). Real-time computing without stable states: a new framework for neural computation based on perturbations. *Neural Comput.* 14, 2531–2560.
- Machens, C.K. (2010). Demixing population activity in higher cortical areas. *Front. Comput. Neurosci.* 4, 126.
- Machens, C.K., Romo, R., and Brody, C.D. (2005). Flexible control of mutual inhibition: a neural model of two-interval discrimination. *Science* 307, 1121–1124.
- Machens, C.K., Romo, R., and Brody, C.D. (2010). Functional, but not anatomical, separation of “what” and “when” in prefrontal cortex. *J. Neurosci.* 30, 350–360.
- Mante, V., Sussillo, D., Shenoy, K.V., and Newsome, W.T. (2013). Context-dependent computation by recurrent dynamics in prefrontal cortex. *Nature* 503, 78–84.
- Martens, J., and Sutskever, I. (2011). Learning recurrent neural networks with Hessian-free optimization. In Gettkr, L. and Scheffer, T., eds., *Proc. 28th Int. Conf. on Machine Learning (ICML-11)*, 1033–1040.
- Maslov, S., and Sneppen, K. (2002). Specificity and stability in topology of protein networks. *Science* 296, 910–913.
- Merchant, H., Crowe, D.A., Robertson, M.S., Fortes, A.F., and Georgopoulos, A.P. (2011). Top-down spatial categorization signal from prefrontal to posterior parietal cortex in the primate. *Front. Syst. Neurosci.* 5, 69.
- Meyers, E.M., Freedman, D.J., Kreiman, G., Miller, E.K., and Poggio, T. (2008). Dynamic population coding of category information in inferior temporal and prefrontal cortex. *J. Neurophysiol.* 100, 1407–1419.
- Miller, E.K., Freedman, D.J., and Wallis, J.D. (2002). The prefrontal cortex: categories, concepts and cognition. *Philos. Trans. R. Soc. Lond. B Biol. Sci.* 357, 1123–1136.
- Miller, E.K., Nieder, A., Freedman, D.J., and Wallis, J.D. (2003). Neural correlates of categories and concepts. *Curr. Opin. Neurobiol.* 13, 198–203.

- Murakami, M., and Mainen, Z.F. (2015). Preparing and selecting actions with neural populations: toward cortical circuit mechanisms. *Curr. Opin. Neurobiol.* 33, 40–46.
- Murray, J.D., Bernacchia, A., Freedman, D.J., Romo, R., Wallis, J.D., Cai, X., Padoa-Schioppa, C., Pasternak, T., Seo, H., Lee, D., and Wang, X.-J. (2014). A hierarchy of intrinsic timescales across primate cortex. *Nat. Neurosci.* 17, 1661–1663.
- Murray, J., Bernacchia, A., Roy, N., Constantinidis, C., Romo, R., and Wang, X.-J. (2017). Stable subspace coding for working memory coexists with heterogeneous neural dynamics in prefrontal cortex. *Proc. Natl. Acad. Sci. USA* 141, 394–399.
- Park, I.M., Meister, M.L., Huk, A.C., and Pillow, J.W. (2014). Encoding and decoding in parietal cortex during sensorimotor decision-making. *Nat. Neurosci.* 17, 1395–1403.
- Rabinovich, M.I., and Varona, P. (2011). Robust transient dynamics and brain functions. *Front. Comput. Neurosci.* 5, 24.
- Rabinovich, M., Volkovskii, A., Lecanda, P., Huerta, R., Abarbanel, H.D., and Laurent, G. (2001). Dynamical encoding by networks of competing neuron groups: winnerless competition. *Phys. Rev. Lett.* 87, 068102.
- Rabinovich, M., Huerta, R., and Laurent, G. (2008). Neuroscience. Transient dynamics for neural processing. *Science* 321, 48–50.
- Rabinovich, M.I., Sokolov, Y., and Kozma, R. (2014). Robust sequential working memory recall in heterogeneous cognitive networks. *Front. Syst. Neurosci.* 8, 220.
- Raposo, D., Kaufman, M.T., and Churchland, A.K. (2014). A category-free neural population supports evolving demands during decision-making. *Nat. Neurosci.* 17, 1784–1792.
- Rigotti, M., Ben Dayan Rubin, D., Wang, X.-J., and Fusi, S. (2010). Internal representation of task rules by recurrent dynamics: the importance of the diversity of neural responses. *Front. Comput. Neurosci.* 4, 24.
- Rigotti, M., Barak, O., Warden, M.R., Wang, X.-J., Daw, N.D., Miller, E.K., and Fusi, S. (2013). The importance of mixed selectivity in complex cognitive tasks. *Nature* 497, 585–590.
- Rishel, C.A., Huang, G., and Freedman, D.J. (2013). Independent category and spatial encoding in parietal cortex. *Neuron* 77, 969–979.
- Roelfsema, P.R., and van Ooyen, A. (2005). Attention-gated reinforcement learning of internal representations for classification. *Neural Comput.* 17, 2176–2214.
- Roelfsema, P.R., Khayat, P.S., and Spekrijse, H. (2003). Subtask sequencing in the primary visual cortex. *Proc. Natl. Acad. Sci. USA* 100, 5467–5472.
- Rombouts, J., Roelfsema, P., and Bohte, S.M. (2012). Neurally plausible reinforcement learning of working memory tasks. In *Advances in Neural Information Processing Systems*, 1871–1879.
- Roy, J.E., Riesenhuber, M., Poggio, T., and Miller, E.K. (2010). Prefrontal cortex activity during flexible categorization. *J. Neurosci.* 30, 8519–8528.
- Rubinov, M., and Sporns, O. (2010). Complex network measures of brain connectivity: uses and interpretations. *Neuroimage* 52, 1059–1069.
- Sarma, A., Masse, N.Y., Wang, X.-J., and Freedman, D.J. (2016). Task-specific versus generalized mnemonic representations in parietal and prefrontal cortices. *Nat. Neurosci.* 19, 143–149.
- Savin, C., and Triesch, J. (2014). Emergence of task-dependent representations in working memory circuits. *Front. Comput. Neurosci.* 8, 57.
- Shafi, M., Zhou, Y., Quintana, J., Chow, C., Fuster, J., and Bodner, M. (2007). Variability in neuronal activity in primate cortex during working memory tasks. *Neuroscience* 146, 1082–1108.
- Singh, R., and Eliasmith, C. (2006). Higher-dimensional neurons explain the tuning and dynamics of working memory cells. *J. Neurosci.* 26, 3667–3678.
- Song, S., Sjöström, P.J., Reigl, M., Nelson, S., and Chklovskii, D.B. (2005). Highly nonrandom features of synaptic connectivity in local cortical circuits. *PLoS Biol.* 3, e68.
- Stokes, M.G., Kusunoki, M., Sigala, N., Nili, H., Gaffan, D., and Duncan, J. (2013). Dynamic coding for cognitive control in prefrontal cortex. *Neuron* 78, 364–375.
- Sussillo, D. (2014). Neural circuits as computational dynamical systems. *Curr. Opin. Neurobiol.* 25, 156–163.
- Sussillo, D., and Barak, O. (2013). Opening the black box: low-dimensional dynamics in high-dimensional recurrent neural networks. *Neural Comput.* 25, 626–649.
- Swaminathan, S.K., and Freedman, D.J. (2012). Preferential encoding of visual categories in parietal cortex compared with prefrontal cortex. *Nat. Neurosci.* 15, 315–320.
- Swaminathan, S.K., Masse, N.Y., and Freedman, D.J. (2013). A comparison of lateral and medial intraparietal areas during a visual categorization task. *J. Neurosci.* 33, 13157–13170.
- Wallis, J.D., Anderson, K.C., and Miller, E.K. (2001). Single neurons in prefrontal cortex encode abstract rules. *Nature* 411, 953–956.
- Wang, X.-J. (2001). Synaptic reverberation underlying mnemonic persistent activity. *Trends Neurosci.* 24, 455–463.
- Wang, X.-J. (2002). Probabilistic decision making by slow reverberation in cortical circuits. *Neuron* 36, 955–968.
- Wohrer, A., Humphries, M.D., and Machens, C.K. (2013). Population-wide distributions of neural activity during perceptual decision-making. *Prog. Neurobiol.* 103, 156–193.
- Wong, K.-F., and Wang, X.-J. (2006). A recurrent network mechanism of time integration in perceptual decisions. *J. Neurosci.* 26, 1314–1328.

STAR★METHODS

KEY RESOURCES TABLE

REAGENT or RESOURCE	SOURCE	IDENTIFIER
Software and Algorithms		
Hessian-Free algorithm	Martens and Sutskever, 2011	http://www.cs.toronto.edu/~ilya/pubs/
DPCA algorithm	Machens et al., 2010; Brendel et al., 2011	https://github.com/machenslab/dPCA

CONTACT FOR REAGENT AND RESOURCE SHARING

Further information and requests for resources should be directed to and will be fulfilled by the Lead Contact, Xiao-Jing Wang (xjwang@nyu.edu).

EXPERIMENTAL MODEL AND SUBJECT DETAILS

All neural data were from Freedman and Assad (2006) and Swaminathan and Freedman (2012), where experimental protocol and recording procedures were described in detail. In summary, four male monkeys (*Macaca mulatta*, weighing from 8.0–14 kg) were trained to indicate whether a test stimulus was in the same category as a previously presented sample stimulus by releasing a lever. Stimuli were high-contrast, 9.0° diameter random dot movies, which moved at 12° per s with 100% coherence. In the first experiment (Freedman and Assad, 2006), the stimulus set comprised 12 motion directions (30° apart) and neural responses were recorded from 156 lateral intraparietal (LIP) neurons in two monkeys (Figure 1B, left panel). In the second study (Swaminathan and Freedman, 2012), 6 motion directions were used (60° apart), and neural data were collected from 76 LIP and 447 prefrontal neurons (PFC) in two other monkeys (Figure 1B, right). To avoid confusion, we refer to the LIP populations from the first and second experiments as the LIP1 and LIP2 datasets respectively. In the first experiment, four test stimuli closest (15°) to the category boundary were removed before recording from monkey H. The second experiment incorporated stimuli at category boundary, which were removed from this analysis. To combine data from four monkeys, all stimulus directions were rotated so that the category boundary corresponds to 0° – 180° axis.

The monkeys were implanted with a head post and recording chamber(s), at coordinates determined by magnetic resonance imaging. In the first experiment, the chamber was centered approximately at P3, L10 to allow access to both the intraparietal sulcus (IPS) and superior temporal sulcus by means of a dorsal approach. The recording chamber was centered approximately 3.0 mm posterior to the intraural line, and 10.0 mm lateral from the midline. For the second experiment, PFC chambers were centered on the principal sulcus and anterior to the arcuate sulcus at ~ 27.0 mm anterior to the intra-aural line, while LIP chambers were positioned over the intraparietal sulcus (IPS) centered ~ 3.0 mm posterior to the intra-aural line. The first study was conducted at Harvard University and all experimental procedures followed Harvard Medical School and National Institutes of Health guidelines. The second study was conducted at University of Chicago, where all procedures were in accordance with the University of Chicago's Animal Care and Use Committee and US National Institutes of Health guidelines.

METHOD DETAILS

Model Architecture and Training

We trained a recurrent network model to perform DMC task. The single-unit dynamics is governed by the following equation:

$$\tau \dot{\mathbf{x}}(t) = -\mathbf{x}(t) + \mathbf{W}_{rr}\mathbf{r}(t) + \mathbf{W}_{ur}\mathbf{u}(t) + \boldsymbol{\eta}(t),$$

where $x_i(t)$ is a synaptic current variable of neuron i at time t and neural activity (firing rate, r_i) is a rectified nonlinear function of x : $r_i = \tanh_+(x_i)$. This constrained firing rates to be positive. The recurrent network has 150 units, and the connectivity matrix (\mathbf{W}_{rr}) is initialized to have 10% probability of connections, where non-zero weights are drawn from a normal distribution of zero mean and SD = 0.28. The neural time constant (τ) is 100 ms. Each recurrent neuron receives an independent white noise input ($\boldsymbol{\eta}(t)$) with zero mean and the final $\sigma_r = 0.6$ (see progressive protocol below for σ_r value during training). The input to the network at a given time ($\mathbf{u}(t)$) is fed into the recurrent network through synaptic weight \mathbf{W}_{ur} , which is initialized similarly to \mathbf{W}_{rr} . The model parameters were set to be in the same range as those in Mante et al. (2013). The changes in these parameters do not affect our qualitative findings.

There are 33 input units. The first 32 units have direction-tuned activity with equally spaced preferred directions from 0° to 360°. When a motion stimulus of direction ϕ appears, the mean activity of input units depends of the unit's preferred direction, θ :

$$u(\theta, \phi) = a \exp\left(-(\phi - \theta)^2 / 2\sigma^2\right),$$

with $\sigma = 43.2^\circ$ and $a = 0.8$. The last input unit signals the appearance of fixation dot with a mean activity = 0.05 during the fixation period. At each time step, input units' activity contain contribution from white noise with $\sigma_u = 0.6$. A standard trial consists of 1 s resting period, followed by 500 ms fixation, 650 ms sample stimulus, 1 s delay, 250 ms test stimulus, and 1 s choice period. The total trial duration $T = 4.4$ s. In the model, the duration of test stimulus was clipped to the approximate reaction time of the monkeys at 250 ms (instead of 650 ms in task protocol) to mimic the fact that visual stimulus was removed after monkeys make movement response. Model dynamics are simulated using Euler update with $\Delta t = 10$ ms.

Five output units linearly read out the synaptic currents of the recurrent circuit:

$$\mathbf{y} = \mathbf{W}_{ro}\mathbf{x}.$$

To train the network, we specified the desired target activity for each output unit and iteratively adjusted all plastic synapses to minimize the discrepancy between the readout activity and the target output. [Figure S1](#) illustrates and explains the target outputs of all five units. To assess the performance of the model, we passed activity of action units to a saturating nonlinear function:

$$\tilde{y}_i = 0.5 \tanh(y_i - 3.0).$$

Reaction time is defined as the time window at which \tilde{y} passes a threshold of 0.85 (in arbitrary units). The network responds match (or non-match) if \tilde{y}_1 (or \tilde{y}_2) passes the threshold within 1.2 s after the test stimulus onset and the activity of the opposite choice does not pass the threshold.

All plastic synapses (input, recurrent, and output), as well as the initial conditions of the network activity, $\mathbf{x}(t=0)$, are updated with a supervised training technique called Hessian-free (HF) algorithm ([Martens and Sutskever, 2011](#); [Mante et al., 2013](#)), which is designed to minimize the error defined as:

$$e = \sum_{k=1}^K \sum_{t=0}^T \sum_{i=1}^{N_o} (\hat{y}_i - y_i)^2.$$

The error is effectively the square of the difference between target output, \hat{y}_i , and network output, y_i , summed over all N recurrent neurons from first time step $t = 0$ to the end of the trial $t = T = 4.4$ s and across all K trials in each training batch. HF belongs to the family of truncated Newton methods, which identifies update directions using second-order curvature and combines geometric insight and optimization heuristics to find solution with relatively low computational resource. In addition to standard regularizations in HF method, we imposed L_1 regularization, which simultaneously minimize the L_1 norm of parameters, $\alpha \sum_{j=1}^n |w_j|$. This constrains the algorithm to find sparse synaptic matrix solutions. $\alpha = 0.001$ controls the contribution of the L_1 regularization term on the objective function.

We employed a progressive training protocol, which started with the simplest version of DMC task and gradually increased task difficulty as model's performance passed criteria (C). This yielded an overall higher success rate and faster training. In the first step, the task involved only two stimuli at the middle of categories (effectively a delayed-match-to-sample task) and individual neurons receive no independent noise (the number of trials per batch, K , is 200; $C = 99\%$). Second, noise is gradually added to the system. We increased units' noise, σ_r , by 0.05 at training batches where the network performance passed the criterion, $C = 87\%$, until desired level of noise is reached ($\sigma_r = 0.6$). Third, synapses with near-zero weight are gradually removed, whereby 5% of smallest synapses are set to zeros and their future updates are set to zeros on training batches with performance larger than the criteria ($C = 87\%$). Synaptic clipping is repeated until the probability of connections in the recurrent weight matrix equals 12%. Finally, the number of stimuli (as well as the number of training trials per iteration) was progressively increased, until the network can perform DMC with 12 motion directions ($K = 2,160$, $C = 87\%$).

Model for Multi-scheme Categorization Tasks

The training method for multi-scheme categorization tasks is identical to the standard DMC task except for the structure of inputs. In the independent-input paradigm, the input population consists of 65 units. One unit encodes the fixation dot and 32 units encode motion directions as described for standard DMC. Another set of 32 units encode the angular locations of a dot stimulus, modeled with periodic Gaussian current profiles as described in Model architecture and training. The outputs for match and non-match trials are identical to the standard task, but provide appropriate match or non-match answers corresponding to the new task rule. In the shared-input paradigm, another 500 ms task epoch was added before the fixation period to provide categorization rule signal. The input population consists of 35 units. The first 33 units are identical to standard DMC. The two additional units represent the task rules (horizontal or vertical categorization boundaries) by a pulse current with a magnitude = 0.3 when the corresponding rule is active and a magnitude of zero otherwise.

Analysis of Model Dynamics and Connectivity

We trained 10 instances of the model and performed the same analyses on them. The results shown reflect behaviors observed across all network realizations.

Noiseless trajectories are simulated by setting σ_r (noise in the firing rate) and σ_u (noise in the input current) to zeros (see [Model Architecture and Training](#) section for definition of σ_r and σ_u). Neural states in [Figures 5](#) and [6](#) are defined by the synaptic currents, \mathbf{x} , at each time point to allow both excitatory and inhibitory (subthreshold) dynamics to be observed; neural states defined by firing rates yield a similar picture. The stable resting state is defined as the terminal steady state when network is simulated with the initial

condition obtained from training, $\mathbf{x}(t=0)$, without noise and input. Stable states associated to choices are determined by running dynamics to terminal states for match or non-match trials without noise or input. All stable states are confirmed to have zero velocity and the eigendecomposition of linearized dynamics around these locations yields only negative eigenvalues, indicating attracting states. Velocity vector fields were computed from the dynamic equation in [Model Architecture and Training](#). The magnitude of velocity is defined as the norm of velocity vectors, $\|\dot{\mathbf{x}}\|^2$. All network trajectories and vector field plots reflect the first two to three largest principal components of all the data in the graph, except in [Figure 4](#) where the procedures mirror the analysis of neural data. The locations of saddle points are determined by the optimization methods in [Sussillo and Barak \(2013\)](#) and [Mante et al. \(2013\)](#).

To compare the trained networks with random networks, we trained six randomly connected networks (RCNs) of N units, whose $N \times N$ synaptic matrix has n non-zero elements per row on average. The non-zero synaptic weights are randomly drawn from a Gaussian distribution of zero mean with variance $1/n$ ($N = 1,500$, $n = 100$) [Barak et al., \(2013\)](#). The input population (33 units) is structured as in the standard network. Approximately 30% of RCN units receive currents from one of the input units. On each trial, the initial synaptic variable, x_i , was drawn independently from a Gaussian distribution to generate variability in neural response. Firing rate is defined as $r_i = \tanh(x_i)$. We simulated approximately 17,000 trials of DMC task and collected average neural activity during the time window of 250 – 500 ms after the test stimulus onset. A support vector machine (SVM) is trained with least-squares method to decode match decision from the RCN's activity using half of the data. The model performance is defined as the accuracy of the trained SVM on decoding another half of the data.

We utilized two measures of functional connectivity, neural response similarity and noise correlation. Neural response similarity (r_{response}) quantifies the similarity or difference in neural encoding during task performance. For example, neurons that strongly prefer red category are very similar to each other ($r_{\text{response}} > 0$), but not similar to neurons that are weakly selective to categories ($r_{\text{response}} \sim 0$) and markedly different from neurons that prefer blue category ($r_{\text{response}} < 0$). r_{response} is defined by the covariance (un-normalized measure of correlation) between synaptic currents of a neural pair under noiseless condition. Note that this is different from r_{signal} in other literature, which is defined as a Pearson correlation between averaged firing rates of a neural pair across stimuli ([Cohen and Kohn, 2011](#)). Noise correlation, r_{noise} , is the Pearson correlation between the rate fluctuations of a neuron pair, averaged across all stimulus directions. Both r_{response} and r_{noise} are functions of time. We calculated them at every 250 ms time window, from the beginning of the sample epoch to the end of test period. Results shown generally hold for all time windows. The range of r_{response} reported are averaged across all time windows and across all model realizations.

For [Figure 7D](#), we also calculated the mean category selectivity duration (CSD) of each neural pair. Neurons with no category tuning (CSD = 0) were removed from the analysis. To investigate the dependence of r_{noise} on CSD, neural pairs were segregated into two groups: one group with mean CSD above 90th percentile and another group with CSD below 90th percentile. Then we compared r_{noise} between the two groups with t test. The same procedure was performed using the synaptic coupling, $c_{ij} = w_{ij} + w_{ji}$, between two neurons instead of r_{noise} and the same results were obtained. Also, since r_{noise} and average neural activity are correlated and this correlation can drive the result, we controlled for average neural activity using ANCOVA. The mean r_{noise} of the two groups is adjusted based on the fitted correlation between r_{noise} and average neural activity across all task conditions before the difference between groups is assessed.

QUANTIFICATION AND STATISTICAL ANALYSIS

Analyzing Temporal Properties of Selectivity

Neurons with average firing rates less than 1 Hz in each and every epoch were removed from all analyses (leaving $N = 156, 74, 380$ for LIP1, LIP2, and PFC, respectively).

Spike trains of individual trials were filtered with a sliding Gaussian kernel of 90 ms width and 30 ms time step. We focused our analysis on the time period from the onset of the fixation epoch to 250 ms after the test stimulus onset. Only correct trials are included in the analysis.

To analyze selectivity time course, we combined different measures. First, for each time window, we tested whether neurons' firing rates are significantly modulated by stimulus categories (t tests, $P < 0.05$, Bonferroni corrected) and defined "a category selectivity phase" as a series of consecutive time windows where neurons are significantly selective to categories. Any selectivity phases shorter than 150 ms were removed from the analysis. The category selectivity duration (CSD) is defined as the duration of its longest category selectivity phase. Categories selectivity phases are plotted in [Figures 3A–3D](#), [S2B](#), and [S2C](#). CSD appears in [Figure 7D](#) and [S2D–S2G](#).

In addition, we used "stimulus selectivity" measure that is very similar to "category selectivity" measure described above. The only difference is for each time window we tested whether firing rates depend significantly on stimulus directions (one-way ANOVA, $P < 0.05$, Bonferroni corrected) rather than stimulus categories. Then we defined "stimulus selectivity phase" as a series of consecutive time windows where neurons are significantly selective to directions. Stimulus selectivity duration is plotted in [Figures S6C–S6E](#).

For category selectivity, we also calculated category selectivity preferences and magnitude at each time window. To achieve this, we assessed the strength of category tuning with a d' measure for each time window in category selectivity phase, defined as:

$$d' = \frac{\mu_1 - \mu_2}{\sqrt{\frac{1}{2}(\sigma_1^2 + \sigma_2^2)}}$$

where μ_i and σ_i is the mean and standard deviation of firing rates in response to stimuli in category i . d' is unbounded and varies from positive values (neuron preferring category blue) to negative values (preferring category red). d' measure is color-coded in [Figures 3A–3D](#).

Neurons are classified into three classes based on the properties of their category selectivity phases. (1) Persistent selectivity describes neurons that have only one selectivity phase, which overlaps sample, delay, and test epochs for at least 150 ms. (2) Partial selectivity refers to neurons that are selective to only one category, but not included in the persistent selectivity group. (3) Switching selectivity covers neurons that switch their category preference at least once during the trial. We calculated the proportion of neurons in each group relative to the total number of category-selective neurons. This classification method is used in [Figure S2A](#).

Another measure of category tuning strength we used is category tuning index (CTI). We identified groups of stimulus pairs with the same distance (direction pairs of 30°, 60°, 90°, 120° apart) and within each group split the pairs into two subgroups (same category v.s. different categories). The CTI measured the difference in firing rate (averaged across all trials for each direction) for each neuron between pairs of directions in different categories (a between category difference) and the difference in activity between pairs of directions in the same category (a within category difference). The CTI was defined as the difference between the within category and between category differences divided by their sum. Values of the index could vary from 1 (strong differences in activity to directions in the two categories) to -1 (large activity differences between directions in the same category, no difference between categories). A CTI value of 0 indicates the same difference in firing rate between and within categories.

Population Response Analysis

To extract population response patterns, we applied demixing principal component analysis (DPCA) to the firing rate traces averaged by task conditions. The algorithmic details and mathematical justification are outlined in [Brendel et al. \(2011\)](#), [Machens et al. \(2005, 2010\)](#), and [Machens \(2010\)](#). In brief, DPCA computed marginalized covariance matrices, denoted C_ϕ , that account for neural response variance due to a subset of task variables, $\phi \in \{t, \theta, \{t, \theta\}\}$. Matrices C_ϕ can be computed by first calculating the marginalized average:

$$\begin{aligned}\bar{y}_t &= \langle \mathbf{r}(t, \theta) \rangle_\theta, \\ \bar{y}_\theta &= \langle \mathbf{r}(t, \theta) \rangle_t, \\ \bar{y}_{t,\theta} &= \mathbf{r}(t, \theta) - \bar{y}_t - \bar{y}_\theta,\end{aligned}$$

and finding marginalized covariance through equation:

$$C_\phi = \langle \bar{y}_\phi \bar{y}_\phi^T \rangle$$

The variance captured by any subset of weight vectors, $\mathbf{W} = [\mathbf{w}_1, \dots, \mathbf{w}_j, \dots]$, due to a variable subspace ϕ is denoted by $v_\phi^2(\mathbf{W}) = \sum_j \mathbf{w}_j^T C_\phi \mathbf{w}_j$. DPCA algorithm searches for a set of orthogonal weight vectors, \mathbf{W}^* , that maximizes the cost function, $L = (\sum_\phi v_\phi^2 / \sum_\phi v_\phi)$. Maximizing L will optimize the trade-off between two objectives: enlarging the numerator that represents the overall variance captured by \mathbf{W}^* and downsizing the denominator to ensure that each component only captures variance in a single variable subspace ([Brendel et al., 2011](#)).

Average neural activity was divided into three portions based on the task epochs: sample (−100–650 ms relative to sample onset), delay (800–1,550 ms), and test (1,600–2,150 ms). For the model, the activity during test epoch came from a slightly wider time window (1,600–2,650 ms) to capture the convergence to match and non-match states, which occurs slightly slower in the model. DPCA analysis was performed on simulated activity of the model in presence of noise, using comparable number of trials as in the data (~400 trials). We first denoised the neural responses by applying regular principal component analysis and focusing on the subspace of M largest components that explain 95% of variance in the data (M may vary from dataset to dataset). Then the activity was passed through the DPCA algorithm, which yields M components ranging from the most stimulus-dependent component (\mathbf{w}_θ), to the component that captures most combined variance of stimulus and time ($\mathbf{w}_{\theta,t}$), to the most time-varying component (\mathbf{w}_t). For the sample and test period, we removed \mathbf{w}_t from the \mathbf{W}^* matrix and derived two-dimensional representation of activity within the remaining subspace using classical multi-dimensional scaling (MDS) ([Borg and Groenen, 1997](#)). Stimuli are represented as vectors in an $M - 1$ dimensional space; each dimension corresponds to a principal component. The MDS algorithm searches for 2D coordinates of stimuli that preserve their pairwise Euclidean distances. For the delay period, we projected activity during −250–2,000 ms onto \mathbf{W}^* , $\mathbf{z}(t) = \mathbf{W}^* \mathbf{Y}(t)$ to visualize the overall neural activity within a subspace spanned by delay-related components. [Figures 4E–4H](#) plotted projected activity on \mathbf{w}_θ and $\mathbf{w}_{\theta,t}$.

Note that since we do not have neurophysiological recordings from one monkey for trials where test stimuli are 15° away from the boundary (see [EXPERIMENTAL MODEL AND SUBJECT DETAILS](#)), we need to remove all trials with test stimuli 15° away from the boundary to construct population trajectories shown in [Figure 4I](#). This step was not applied to other analyses in this paper.

DATA AND SOFTWARE AVAILABILITY

Software for modeling and data analysis is written in MATLAB. Requests for source code and data should be directed to our Lead Contact.